

Research Paper Theory vs Code Implementation

RESEARCH PAPER: "Optimizing the Utilization of LLMs"

1. PROBLEM STATEMENT

- Multiple LLM models with different costs and capabilities
Example: GPT-4 (\$0.06/1K tokens) vs GPT-3.5 (\$0.002/1K tokens)
- Same task can be done by cheap OR expensive model
- Challenge: Maximize accuracy while minimizing total cost
- Trade-off: Cost ↔ Accuracy

2. PROPOSED SOLUTION: ML-Based Predictor (MLBP)

Step 1: Ground Truth Collection

- Sample N tasks
- Query ALL models for each task
- Record: model_i correctly classified task_j? (YES/NO)

Step 2: Feature Extraction

- Task features: length, complexity, domain
- Model features: parameters, cost, capability

Step 3: Train ML Predictor

- Input: (task, model) → Features
- Output: P(model succeeds on task)
- Algorithm: Random Forest, XGBoost, Neural Net

Step 4: Predict for New Tasks

- Given: new task + all models
- Predict: P(success) for each model
- Result: Probability matrix [N_tasks × N_models]

3. MULTI-OBJECTIVE OPTIMIZATION

Objective 1: Maximize Total Accuracy = $\sum P(\text{model}_i \text{ succeeds on task}_i)$
Objective 2: Minimize Total Cost = $\sum \text{cost}(\text{model}_i, \text{task}_i)$

Constraint: Assign exactly ONE model to each task

Multi-Objective Algorithms:

- NSGA-II (Non-dominated Sorting Genetic Algorithm)
- SPEA2 (Strength Pareto Evolutionary Algorithm)
- Random Search (baseline)

Result: Pareto Front

- Set of non-dominated solutions
- Each point: (cost, accuracy) pair
- User picks based on budget/accuracy requirement

4. EVALUATION METRICS

- IGD (Inverted Generational Distance): How close to reference Pareto front?
Lower is better. 0 = perfect match.
- Δ (Delta): How evenly spread are solutions on Pareto front?
Lower is better. Good coverage of tradeoffs.
- M_n (Number of solutions): How many options does user have?
Higher is better. More choices.

5. BASELINE STRATEGIES

- Cheapest Model Only: Use GPT-3.5 for everything (low cost, low accuracy)
- Best Model Only: Use GPT-4 for everything (high cost, high accuracy)
- Random Assignment: Randomly assign models (baseline)
- Compare optimizers against these baselines

CODE IMPLEMENTATION (My Project)

1. PROBLEM ADAPTATION

- Instead of multiple MODELS → Use multiple PROMPTS simple (10 tokens), standard (25), fewshot_1 (50), fewshot_3 (90)
- Same LLM (Ollama gemma3) with different instruction complexity
- Dataset: HDFS logs (5,000 jobs) with anomaly labels
- Trade-off: Cost (tokens) ↔ Accuracy (detection rate)

2. MLBP IMPLEMENTATION: predictors/mlbp.py

```
create_training_data(jobs, prompts):
    • Loop: 5000 jobs × 4 prompts = 20,000 pairs
    • For each pair: Call query_gemini(job, prompt) → LLM prediction
    • Compare with ground_truth → label = 1 if correct, 0 if wrong
    • Cache results to disk (results/ollama_cache/{hash}.json)
```

```
Feature Engineering (7 features):
    • job["tokens"], job["unique_tokens"], job["has_error"]
    • prompt["tokens"], prompt["id"]
    • job_tokens/prompt_tokens, job_tokens*prompt_tokens
```

```
train_predictor(X, y):
    • Algorithm: XGBoost (400 estimators, depth=8, lr=0.05)
    • Input: X (20000×7), y (20000×1)
    • Output: Trained model
    • Current accuracy: 7%, Target: 85%
```

```
predict_probabilities(model, jobs, prompts):
    • Predict for all (job, prompt) pairs
    • Result: P(success) matrix [5000 × 4]
    • Used by optimizers to make decisions
```

3. MULTI-OBJECTIVE OPTIMIZATION: optimizers/

Maximize: total_accuracy = $\sum P(\text{prompt}_i \text{ succeeds on job}_i)$
Minimize: total_cost = $\sum \text{prompt}_i \text{tokens}$

Constraint: Each job assigned exactly one prompt

Implemented Algorithms:
 • random_search.py: 1000 random assignments
 • nsga2.py: Pop=100, Gen=50 (BEST: IGD=0.0000)
 • spea2.py: Archive-based, fitness sharing

```
evaluation/pareto.py:
    • Extract non-dominated solutions
    • Pareto front: ~40 solutions for NSGA-II
    • visualization/plots.py: Generate plots
```

4. EVALUATION: evaluation/metrics.py

```
compute_igd(pf, reference): Distance to reference Pareto front
NSGA-II: 0.0000 ✓ (perfect)
```

```
compute_delta(pf): Spread of solutions
Lower = better diversity
```

```
compute_mn(pf): Count solutions on Pareto front
NSGA-II: ~40 solutions
```

5. BASELINE: evaluation/baseline.py

```
single_prompt_baseline(jobs, prompt, probabilities):
    • Test 4 strategies: simple/standard/fewshot_1/fewshot_3 for ALL jobs
    • Results: simple: 75.1%, standard: 77.7%, fewshot_1: 79.0%, fewshot_3: 84.1%
    • Saved to: results/tables/baseline_results.csv
```

Maps to

Implements

Same algorithms

Same metrics

KEY DIFFERENCES: Paper vs Implementation

Aspect | Paper Theory | Your Implementation

Models/Prompts | Multiple LLMs | Single LLM, 4 prompts

Dataset | General NLP tasks | HDFS anomaly detection

LLM API | OpenAI GPT-3.5/4 | Local Ollama (gemma3)

Cost measure | \$ per 1K tokens | Token count (proxy)